

Iterators:

- Many operations access only one tuple at a time.
 - read tuple.
 - inspect
 - dispose
 - read next tuple...

Open() — initiates the process

GetNext() — return next tuple

close() — ends process

Example:

$\pi_a \sigma_{b=3} R$

π_a

|

on the fly

$\sigma_{b=3}$

|

seq scan of R

R

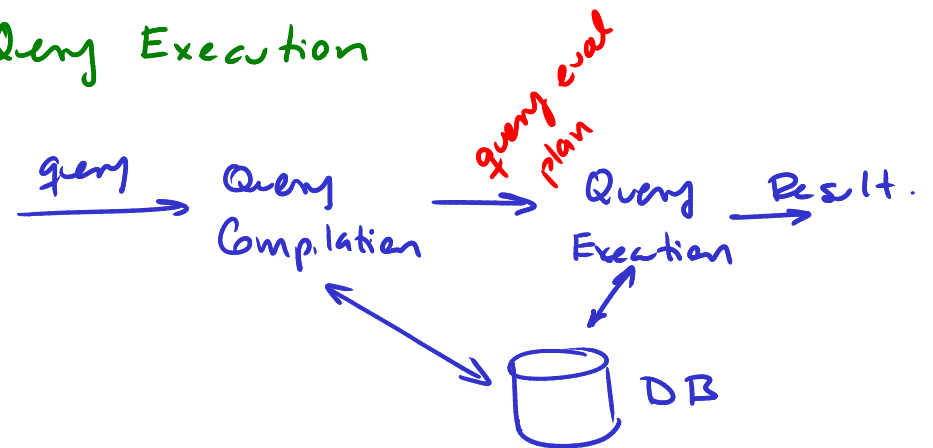
π and σ can be implemented as iterators

σ inspects one tuple at a time, sends one tuple at a time to π

No need to store any tuple in memory

④

Query Execution



Query Compilation

- a) Parsing. A parse tree is constructed
- Create an algebraic expression.

- b) Query Rewrite:

- Several equivalent query expression

- c) Physical plan generation

- Each expression is converted to an evaluation plan by indicating the alg. to use.

b) and c) are the **query optimizer**
⇒ find best query plan.



①

- 1) Which algebraic expression is the one leading to the most efficient alg.
- 2) For each operation in the expression which alg. will be used to answer it.
- 3) How should each operation pass data to the next operation.
- 4) How are the relations going to be accessed.

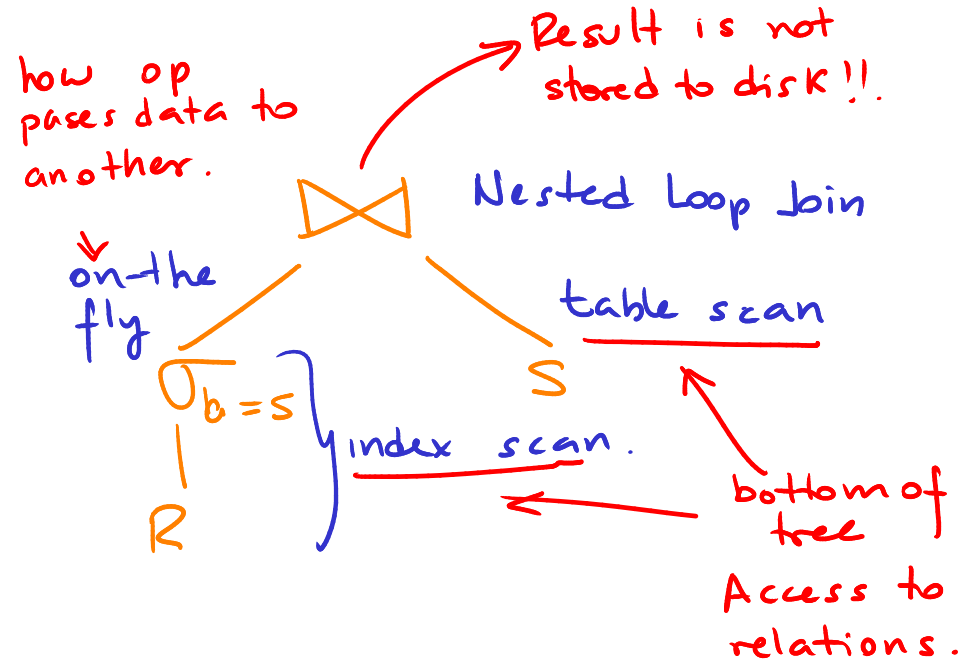
Ex: $R(a, b)$ $S(a, c)$

SELECT * from R natural join S
WHERE $b = 5$

Equivalent Expressions



Annotate tree with algorithms and access methods.



Estimate cost.

\Rightarrow choose fastest!

Access to tuples:

- Sequential scan of heap of Rel.
- or
- Using an index to scan a subset of tuples of R (index scan)

Result of query:

- Kept in memory.