

U,  $\cap$ ,  $-$

We can also use TPMM S

- Do phase one of R
- Do phase one of S.
- Phase 2: do both R and S at the same time:

Read 1 block of each section of R and S at a time:

$\Rightarrow$  do operation on tuples in memory.

$$\text{We need } \left\lceil \frac{B(R)}{M} \right\rceil + \left\lceil \frac{B(S)}{M} \right\rceil < M$$

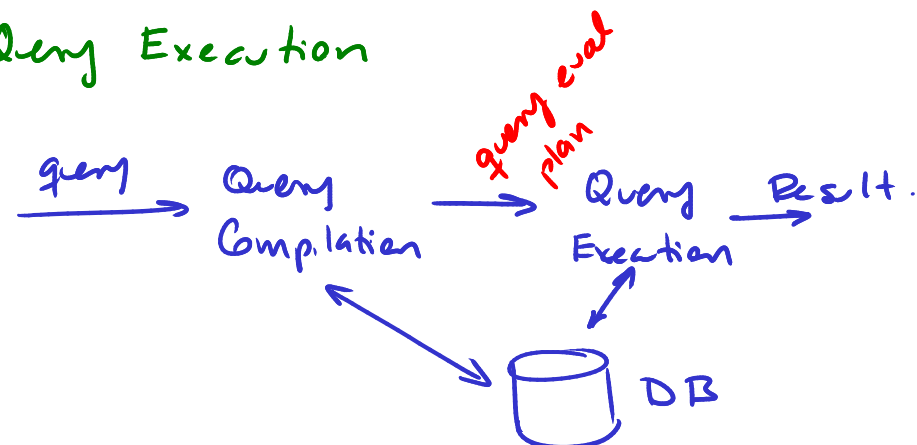
for second pass:

$\Rightarrow$  Memory required is approx.

$$B(R) + B(S) \leq M^2$$

Cost:  $3(B(R) + B(S))$

## Query Execution



## Query Compilation

a) Parsing. A parse tree is constructed

- Create an algebraic expression.

b) Query Rewrite:

- Several equivalent query expression

c) Physical plan generation

- Each expression is converted to an evaluation plan by indicating the alg. to use.

b) and c) are the **query optimizer**  
 $\Rightarrow$  find best query plan.

- 1) Which algebraic expression is the one leading to the most efficient alg.
- 2) For each operation in the expression which alg. will be used to answer it.
- 3) How should each operation pass data to the next operation.
- 4) How are the relations going to be accessed.

Ex:  $R(a, b)$   $S(a, c)$

SELECT \* from R natural join S  
WHERE  $b = 5$

Equivalent Expressions



## Duplicate elimination $\delta(R)$

- Sort R using TPMMS
- During second phase, output only first tuple of each set of duplicates

Mem required:

$$B(R) \leq M^2$$

Cost:  $3B(R)$

## Group By $\gamma$

Use TPMMS to sort by aggr. attributes  
Like  $\delta(R)$ , during second phase  
for each group of tuples in output  
compute aggregation  
output result

Requires one pass of tuples in group.  
Memory required for computing agg. is  
less than 1 block.

Total mem required  $B(R) \leq M^2$

Cost:  $3B(R)$

Memory required

$$\left\lceil \frac{B(R)}{M} \right\rceil \leq M - 1$$

$\Rightarrow$  Approximately  $B(R) \leq M^2$

Cost:

Phase 1:  $B(R)$  Read  $B(R)$  Write.

Phase 2:  $B(R)$  Read

Assume Read = Write

$\Rightarrow 3 B(R)$

100.

100.

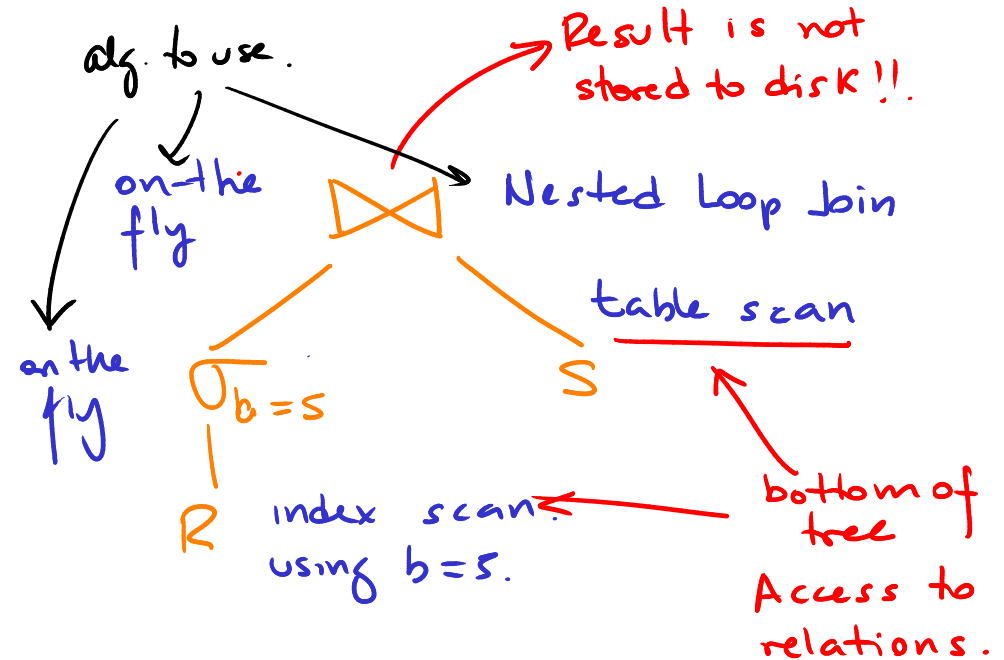
and output is sorted.

We can generalize # passes to.

$$\left\lceil \log_{M-1} B(R) \right\rceil \dots$$

But usually with a decent amount of memory we can sort very large relations in 2 passes.

Annotate tree with algorithms and access methods.



Estimate cost.

$\Rightarrow$  choose fastest!

Access to tuples:

- Sequential scan of heap of Rel.

or

- Using an index to scan a subset of tuples of  $R$  (index scan)

Result of query:

- Kept in memory.

Iterators:

- Many operations access only one tuple at a time.
  - read tuple.
  - inspect
  - dispose
  - read next tuple...

Open() — initiates the process

GetNext() — return next tuple

Close() — ends process

Example:

$\pi_a \sigma_{b=3} R$

$\pi_a$  on the fly

$\sigma_{b=3}$  on the fly

seq scan of  $R$

$\pi$  and  $\sigma$  can be implemented as iterators

$\sigma$  inspects one tuple at a time, sends one tuple at a time to  $\pi$

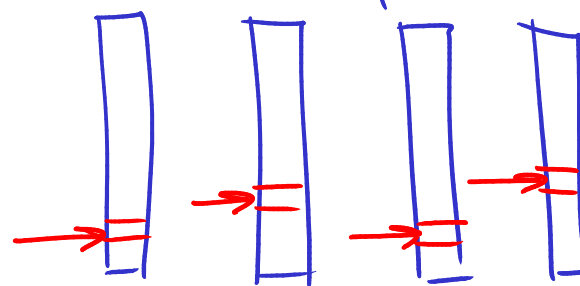
No need to store any tuple in memory

(4)

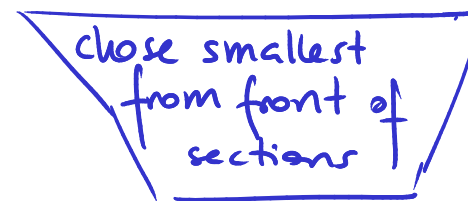
If # sorted sections  $\leq M-1$   
then

Phase 2:

- Merge sorted sections by reading one block of each section at a time.
- Use 1 block for output.



sorted sections of at most  $M-1$  block.



output sorted tuple.

if # sections  $\geq M$  we might need 3 or more phases

## Two pass algorithms based on sorting

Algorithms that read data twice.

- Read tuples
  - Process tuples
  - Write tuple to disk
  - Read tuples
  - Process tuples.
  - ⇒ output result
- } first pass.
- } second pass.

## Sorting $T$

- By sorting we can implement other operations (eg.  $\cap$ ,  $\cup$ ,  $\bowtie$ ).

## Two Phase Multiway Merge Sort **TPMMS**

- Alg. to sort large relations.

$$B(R) > M$$

• Phase I:

For each  $M$  blocks of  $R$   
Read  $M$  blocks.

Sort

Write back to temp. storage.

This creates  $\left\lceil \frac{B(R)}{M} \right\rceil$  sorted sections

## Parameters to measure cost

$M$ . Amount of memory available  
in number of blocks

$B(R)$  # of blocks used by heap of  $R$

$|R|$  # of tuples of  $R$  (book uses  $T(R)$ )

$V(R, a)$  # of different values of att  $a$   
in  $R$

In general:

$$V(R, [a_1, a_2 \dots a_n])$$

$$= |\gamma_{a_1, a_2 \dots a_n} R|$$

⇒ # of different values for tuple  
 $a_1 \dots a_n$

## Cost Model

- We assume that the major component of cost is I/O
- Cost of read equal to cost of write
- Cost of random access of pages equal to cost of seq access.

## Algorithms to answer queries.

2 main classifications.

a) based on type of algorithm:

- 1) Sorting based
- 2) Hash based
- 3) Index based

b) based on difficulty.

- 1) One-pass: Relations are read only once.
- 2) Two passers.
  - Read data (1st pass)
  - Process.
  - Write data.
  - Read data again. (2nd pass).

2nd pass might read diff number of blocks than 1st pass.

- 3) Three or more passers.  
(needed for very large relations).
  - Generalization of Two passers.

Because tables are usually large we approximate to:

Cost:

$$B(R) \cdot \left\lceil \frac{B(S)}{M} \right\rceil \approx \frac{B(R) \cdot B(S)}{M}.$$

It does not really matter which relation we read in the outside loop. But outside table only read once.

That might affect which table is outside.

## Block based Nested Join.

Generalization of 1 pass join.

- What if no relation fits in memory?

Assume:  $B(S) > M$ .

outside loop.  $B(R) > M$

For each M-1 blocks of S  
Read blocks and organize them in mem.  
For each block of R.  
read tuples in block  
for every tuple r in R  
find matching tuples in  
read blocks of S.

inside loop.

Each block of R is read

$\left\lceil \frac{B(S)}{M} \right\rceil$  times.

We also need to read S.

Cost:

$$\left\lceil \frac{B(S)}{M} \right\rceil \cdot B(R) + B(S)$$

## One Pass Alg.

1) Tuple-at-a time  $\Pi, \sigma$

- We can read one block at a time.  
 $\Rightarrow$  use one memory buffer.

$\Pi_a$   
|  
R

- Read one block at a time,
- Inspect each tuple,  
output result
- Repeat.

or

if we received tuples from another operation, one tuple at a time with no need for buffering.

(on the fly — no memory needed)

$\Pi_a$

on the fly.



- Receive tuple from  $\Pi$  via iterator.
- Output result
- Repeat.

No block in memory needed.

But assume 1 block for simplicity's sake.

Other one pass unary operators.


Duplicate elimination ( $\delta$ )

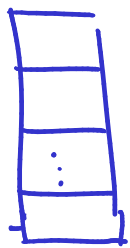
• Read each tuple.

• If we have seen it, ignore

• Otherwise output and keep track of it.

We need to keep a copy of each distinct tuple.

input  
tuples   
(iterator or  
from R heap)

 at most  
 $M-1$   
available  
for  
distinct.

We do not need block for output.

⇒ tuple in result output immediately.

We can do  $\delta R$  in one pass  
as long as:

$$B(\delta(R)) \leq M-1.$$

Book uses:

$$B(\delta(R)) \leq M$$

because  $M \gg 1$

So use latter for consistency.

⑧

Summary of 1 pass algorithm.

	<u>Approx blocks of</u> <u>M required</u>
$\pi, \sigma, \cup_B$	1
$\gamma, \delta$	$B(R)$
$\cup$	$B(R) + B(S)$
$\cap, \cap_B, -$	$\} \min(B(R), B(S))$
$-B, \bowtie, \times$	

order by is a variation of  $\gamma, \delta$   
denoted  $\gamma$



We always read smaller table into M

To compute  $R - S$ ,  $R \bowtie S$ .

Read S  
for every tuple  $t$  in R  
if  $t$  not in S  
output  
(for — also keep track of those output)

To compute  $S - R$ ,  $S \bowtie R$

Read S  
for — remove all duplicates at the same time.  
For every tuple  $t$  in R  
if  $t$  in S  
remove from S  
for — remove one duplicate only  
Output tuples left in S

But, how do we know  $B(\delta(R))$  without calculating  $\delta(R)$  first?

⇒ Statr.

$R(a_1, a_2 \dots a_n)$

then.

We can use  $V(R, a_1 \dots a_n)$  and the size of the tuple in R to calculate  $\delta(R)$ .

Group By:

Generalization of  $\delta(R)$

Remember

$$\delta(R) = \gamma^{a_1 \dots a_n} R$$

For  $\gamma^{<attlist>} R$ .  
 $<explist>$

We need to keep track of:

- Each different value of  $<attlist>$ .
- Info needed to compute  $<explist>$ .

- $\min(x)$   $\left\{ \begin{array}{l} \text{Keep current min/max} \\ \text{max}(x) \end{array} \right.$
- $\text{sum}(x)$  • Keep current sum
- $\text{count}(x)$  Keep current count
- $\text{avg}(x)$  Keep both current count and sum.

We cannot output tuples until we have read all input tuples.

- We must also create access structures in memory (hash tables, b+ trees) to efficiently find group tuple belongs to.

### • In general

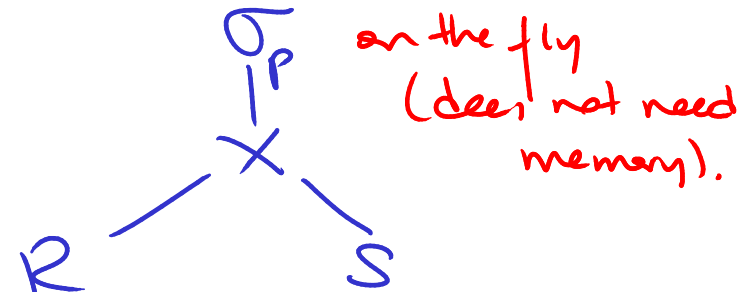
- The amount of memory required per group is small.
- Proportional to the number of different groups.

$$|\chi_{\langle \text{explicit} \rangle}^{a_1 \dots a_i} R| \propto V(R, a_1 \dots a_i)$$

⊗

$$R \bowtie S = \sigma_p(R \times S)$$

Since we can do  $\sigma_p$  on the fly.



But join is common, so DBMS optimize it:

Read S  
 for every tuple  $t$  in R  
 for every tuple  $s$  in S  
 if  $t$  and  $s$  satisfy  $p$   
 output  $\text{join}(t, s)$

— B —

Like  $\cap$ ,  $\cup$ , etc. we load smaller table into memory.

But algorithm is different depending on which table is smaller:

$\cap, \cap_B, \times, \bowtie, -, -B$ .

- All commutative operations.
- Keep smaller table in memory (plus data structures for fast access).
- Plus at most one block for other table:

One pass if, approximately:

$$\min(B(R), B(S)) \leq M.$$

Specifically for each of these operations:

Because they are commutative, assume

$$B(R) \geq B(S)$$

$\cap, \cap_B$

Read  $S$ , organize in data structure.

for every tuple  $t$  in  $R$

if  $t$  in  $S$

if bag op  $\Rightarrow$  output  $t$  if needed

otherwise output  $t$  first time only.

$\times$

Read  $S$

for every tuple  $t$  in  $R$

for every tuple  $s$  in  $S$

compute cross product, output.

(14)

We can do it in one pass if we have enough memory to

- hold all different groups
- data structures for quick access to groups.
- any data required to compute grouping function.

In general size of tuple of result much smaller than original tuple.

So we simplify

We can do group-by in one pass if

$$B(R) \leq M$$

(11)

One Pass alg. for binary operations:

$\cup, \cap, -, \times, \bowtie$

In practice set operations of two types:

- The sets: No duplicates (default).
- Bags: duplicates.

UNION  
INTERSECT  
EXCEPT

} ALL

$\Rightarrow$  Represented  $\cup_B, \cap_B, -_B$

TABLE R UNION ALL TABLE S

Result contains all tuples in R plus all tuples in S.

TABLE R INTERSECT ALL TABLE S

if a tuple in R has  $m$  duplicates in R  
and  $n$  duplicates in S  
result contains  $\min(m, n)$  duplicates  
of tuple.

TABLE R EXCEPT ALL TABLE S

if a tuple in R has  $m$  duplicates in R  
and  $n$  duplicates in S  
result contains  $\min(m - n, 0)$

(12)

$\cup_B$

- Similar to  $\pi$ :
- We only need to inspect one tuple at a time.  
 $M = 1$ , regardless of size of input.

$\cup$

- Removes duplicates:
- Equivalent to  $\delta(R \cup_B S)$

The book is wrong. It states we only need to read S in  $M-1$  and do one-tuple-at-a-time for R (page 716)

We can do in one pass if

$$\delta(R \cup_B S) \leq M$$

We can approximate to:

$$\delta(B(R)) + \delta(B(S)) \leq M$$

We can remove duplicates as we read tuples:

if tuple already read, ignore  
otherwise  $\hookrightarrow$  output  
add to read tuples. (13)