

U, n, -

We can also use TPMM S

- Do phase one of R
- Do phase one of S.
- Phase 2: do both R and S at the same time:

Read 1 block of each section of R and S at a time:

⇒ do operation on tuples in memory.

$$\text{We need } \left\lceil \frac{B(R)}{M} \right\rceil + \left\lceil \frac{B(S)}{M} \right\rceil < M$$

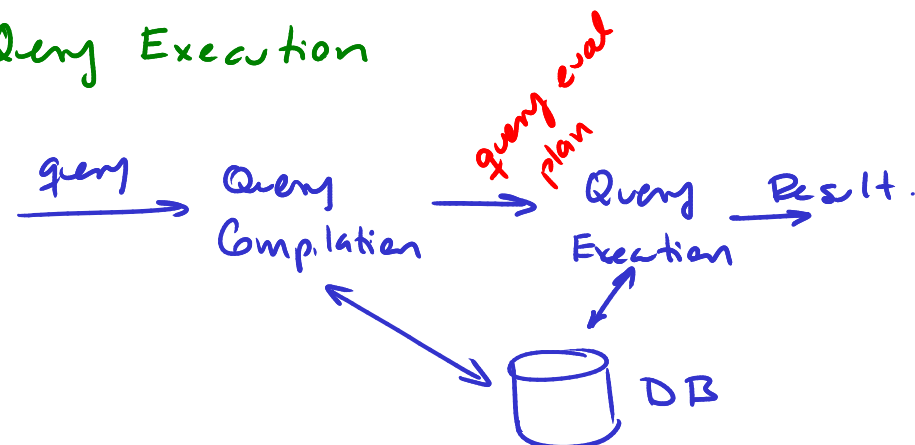
for second pass:

⇒ Memory required is approx.

$$B(R) + B(S) \leq M^2$$

Cost: $3(B(R) + B(S))$

Query Execution



Query Compilation

a) Parsing. A parse tree is constructed

- Create an algebraic expression.

b) Query Rewrite:

- Several equivalent query expression

c) Physical plan generation

- Each expression is converted to an evaluation plan by indicating the alg. to use.

b) and c) are the **query optimizer**
⇒ find best query plan.

- 1) Which algebraic expression is the one leading to the most efficient alg.
- 2) For each operation in the expression which alg. will be used to answer it.
- 3) How should each operation pass data to the next operation.
- 4) How are the relations going to be accessed.

Ex: $R(a, b)$ $S(a, c)$

SELECT * from R natural join S
WHERE $b = 5$

Equivalent Expressions



Duplicate elimination $\delta(R)$

- Sort R using TPMMS
- During second phase, output only first tuple of each set of duplicates

Mem required:

$$B(R) \leq M^2$$

Cost: $3B(R)$

Group By γ

Use TPMMS to sort by aggr. attributes
Like $\delta(R)$, during second phase
for each group of tuples in output
compute aggregation
output result

Requires one pass of tuples in group.
Memory required for computing agg. is
less than 1 block.

Total mem required $B(R) \leq M^2$

Cost: $3B(R)$

Memory required

$$\left\lceil \frac{B(R)}{M} \right\rceil \leq M - 1$$

\Rightarrow Approximately $B(R) \leq M^2$

Cost:

Phase 1: $B(R)$ Read $B(R)$ Write.

Phase 2: $B(R)$ Read

Assume Read = Write

$\Rightarrow 3 B(R)$

100.

100.

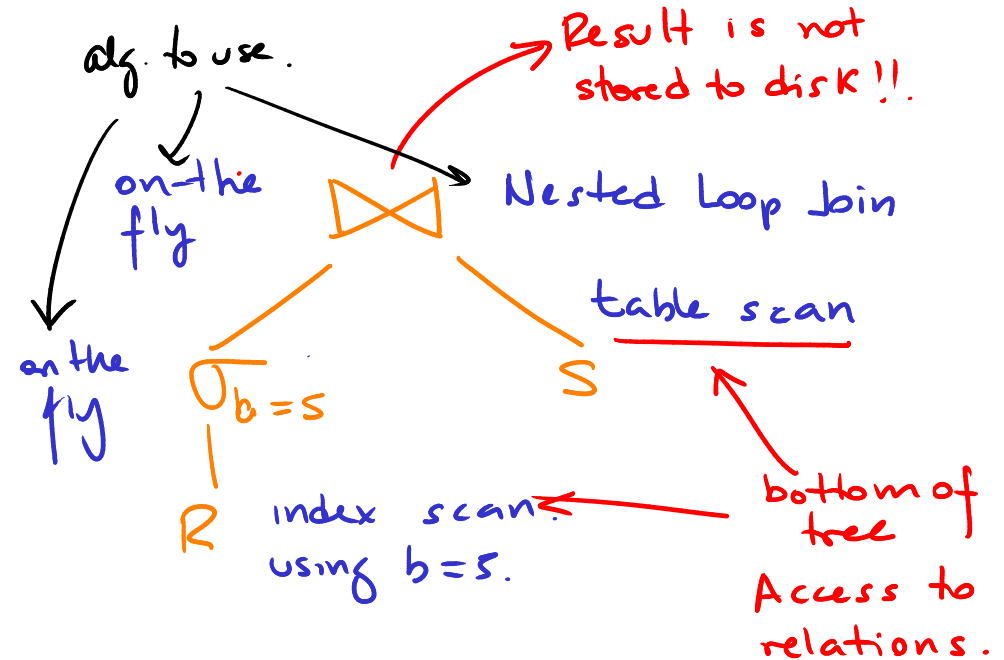
and output is sorted.

We can generalize # passes to.

$$\left\lceil \log_{M-1} B(R) \right\rceil \dots$$

But usually with a decent amount of memory we can sort very large relations in 2 passes.

Annotate tree with algorithms and access methods.



Estimate cost.

\Rightarrow choose fastest!

Access to tuples:

- Sequential scan of heap of Rel.

or

- Using an index to scan a subset of tuples of R (index scan)

Result of query:

- Kept in memory.

Iterators:

- Many operations access only one tuple at a time.
 - read tuple.
 - inspect
 - dispose
 - read next tuple...

Open() — initiates the process

GetNext() — return next tuple

Close() — ends process

Example:

$\pi_a \sigma_{b=3} R$

π_a on the fly

$\sigma_{b=3}$ on the fly

seq scan of R

π and σ can be implemented as iterators

σ inspects one tuple at a time, sends one tuple at a time to π

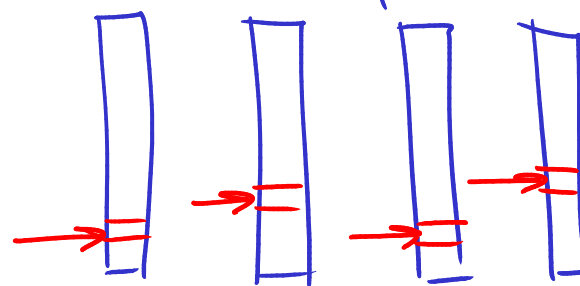
No need to store any tuple in memory

(4)

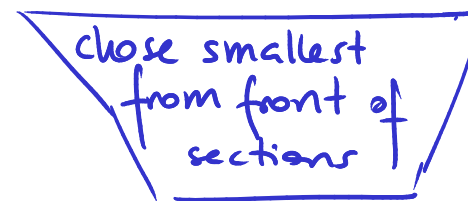
If # sorted sections $\leq M-1$
then

Phase 2:

- Merge sorted sections by reading one block of each section at a time.
- Use 1 block for output.



sorted sections of at most $M-1$ block.



output sorted tuple.

if # sections $\geq M$ we might need 3 or more phases

Two pass algorithms based on sorting

Algorithms that read data twice.

- Read tuples
 - Process tuples
 - Write tuple to disk
 - Read tuples
 - Process tuples.
 - ⇒ output result
- } first pass.
- } second pass.

Sorting T

- By sorting we can implement other operations (eg. \cap , \times , \bowtie).

Two Phase Multiway Merge Sort **TPMMS**

- Alg. to sort large relations.

$$B(R) > M$$

• Phase I:

For each M blocks of R

Read M blocks.

Sort

write back to temp. storage.

This creates $\left\lceil \frac{B(R)}{M} \right\rceil$ sorted sections

Parameters to measure cost

M . Amount of memory available
in number of blocks

$B(R)$ # of blocks used by heap of R

$|R|$ # of tuples of R (book uses $T(R)$)

$V(R, a)$ # of different values of att a
in R

In general:

$$V(R, [a_1, a_2 \dots a_n])$$

$$= |\gamma_{a_1, a_2 \dots a_n} R|$$

⇒ # of different values for tuple
 $a_1 \dots a_n$

Cost Model

- We assume that the major component of cost is I/O
- Cost of read equal to cost of write
- Cost of random access of pages equal to cost of seq access.

Algorithms to answer queries.

2 main classifications.

a) based on type of algorithm:

- 1) Sorting based
- 2) Hash based
- 3) Index based

b) based on difficulty.

- 1) One-pass: Relations are read only once.
- 2) Two passers.
 - Read data (1st pass)
 - Process.
 - Write data.
 - Read data again. (2nd pass).

2nd pass might read diff number of blocks than 1st pass.

- 3) Three or more passers.
(needed for very large relations).
 - Generalization of Two passers.

Because tables are usually large we approximate to:

Cost:

$$B(R) \cdot \left\lceil \frac{B(S)}{M} \right\rceil \approx \frac{B(R) \cdot B(S)}{M}.$$

We should still read smallest table in the outside loop, but that cost might be negligible

Block based Nested Join.

Generalization of 1 pass join.

- What if no relation fits in memory?

Assume: $B(S) > M$.

outside loop.

$B(R) > M$

For each M-1 blocks of S
Read blocks and organize them in mem.
For each block of R.
read tuples in block
for every tuple r in R
find matching tuples in
read blocks of S.

inside loop.

Each block of R is read

$\left\lceil \frac{B(S)}{M} \right\rceil$ times.

We also need to read S.

Cost:

$$\left\lceil \frac{B(S)}{M} \right\rceil \cdot B(R) + B(S)$$

Outside table: smaller one.

One Pass Alg.

1) Tuple-at-a time Π, \Join

- We can read one block at a time.
 \Rightarrow use one memory buffer.

Π_a
|
R

- Read one block at a time,
- Inspect each tuple,
output result
- Repeat.

or

if we received tuples from another operation, one tuple at a time with no need for buffering.

(on the fly — no memory needed)

Π_a
|
 \bowtie

on the fly.

- Receive tuple from R via iterator.
- Output result
- Repeat.

No block in memory needed.

But assume 1 block for simplicity's sake.

Other one pass unary operators.


Duplicate elimination (δ)

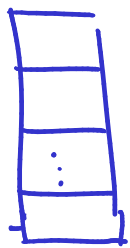
• Read each tuple.

• If we have seen it, ignore

• Otherwise output and keep track of it.

We need to keep a copy of each distinct tuple.

input
tuples 
(iterator or
from R heap)

 at most
 $M-1$
available
for
distinct.

We do not need block for output.

⇒ tuple in result output immediately.

We can do δR in one pass
as long as:

$$B(\delta(R)) \leq M-1.$$

Book uses:

$$B(\delta(R)) \leq M$$

because $M \gg 1$

So use latter for consistency.

⑧

Summary of 1 pass algorithm.

	<u>Approx blocks of</u> <u>M required</u>
π, σ, U_B	1
γ, δ	$B(R)$
U	$B(R) + B(S)$
$\cap, \cap_D, -$	$\} \min(B(R), B(S))$
$-B, \bowtie, \times$	

order by is a variation of γ, δ
denoted γ

We always read smaller table into M

To compute $R - S$, $R \bowtie S$.

Read S
for every tuple t in R
if t not in S
output
(for — also keep track of those output)

To compute $S - R$, $S \bowtie R$

Read S
for — remove all duplicates at the same time.
For every tuple t in R
if t in S
remove from S
for — remove one duplicate only
Output tuples left in S

But, how do we know $B(\delta(R))$ without calculating $\delta(R)$ first?

⇒ Statr.

$R(a_1, a_2 \dots a_n)$

then.

We can use $V(R, a_1 \dots a_n)$ and the size of the tuple in R to calculate $\delta(R)$.

Group By:

Generalization of $\delta(R)$

Remember

$$\delta(R) = \gamma^{a_1 \dots a_n} R$$

For $\gamma^{<attlist>} R$.
 $<explist>$

We need to keep track of:

- Each different value of $<attlist>$.
- Info needed to compute $<explist>$.

- $\min(x)$ $\left\{ \begin{array}{l} \text{Keep current min/max} \\ \text{max}(x) \end{array} \right.$
- $\text{sum}(x)$ • Keep current sum
- $\text{count}(x)$ Keep current count
- $\text{avg}(x)$ Keep both current count and sum.

We cannot output tuples until we have read all input tuples.

- We must also create access structures in memory (hash tables, b+ trees) to efficiently find group tuple belongs to.

• In general

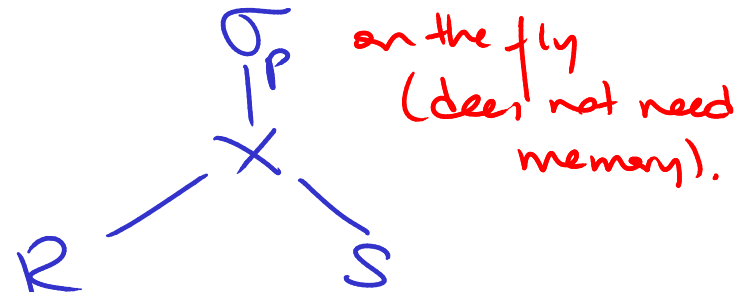
- The amount of memory required per group is small.
- Proportional to the number of different groups.

$$|\chi_{\langle \text{explicit} \rangle}^{a_1 \dots a_i} R| \propto V(R, a_1 \dots a_i)$$

⊗

$$R \bowtie S = \sigma_p (R \times S)$$

Since we can do σ_p on the fly.



But join is common, so DBMS optimize it:

Read S
 for every tuple t in R
 for every tuple s in S
 if t and s satisfy p
 output $\text{join}(t, s)$

— B —

Like \cap , \cup , etc. we load smaller table into memory.

But algorithm is different depending on which table is smaller:

$\cap, \cap_B, \times, \bowtie, -, -B$.

- All commutative operations.
- Keep smaller table in memory (plus data structures for fast access).
- Plus at most one block for other table:

One pass if, approximately:

$$\min(B(R), B(S)) \leq M.$$

Specifically for each of these operations:

Because they are commutative, assume

$$B(R) \geq B(S)$$

\cap, \cap_B

Read S , organize in data structure.

for every tuple t in R

if t in S

if bag op \Rightarrow output t if needed

otherwise output t first time only.

\times

Read S

for every tuple t in R

for every tuple s in S

compute cross product, output.

(14)

We can do it in one pass if we have enough memory to

- hold all different groups
- data structures for quick access to groups.
- any data required to compute grouping function.

In general size of tuple of result much smaller than original tuple.

So we simplify

We can do group-by in one pass if

$$B(R) \leq M$$

(11)

One Pass alg. for binary operations:

$\cup, \cap, -, \times, \bowtie$

In practice set operations of two types:

- The sets: No duplicates (default).
- Bags: duplicates.

UNION
INTERSECT
EXCEPT } ALL

\Rightarrow Represented $\cup_B, \cap_B, -_B$

TABLE R UNION ALL TABLE S

Result contains all tuples in R plus all tuples in S.

TABLE R INTERSECT ALL TABLE S

if a tuple in R has m duplicates in R
and n duplicates in S
result contains $\min(m, n)$ duplicates
of tuple.

TABLE R EXCEPT ALL TABLE S

if a tuple in R has m duplicates in R
and n duplicates in S
result contains $\min(m - n, 0)$

(12)

\cup_B

- Similar to π :
- We only need to inspect one tuple at a time.
 $M = 1$, regardless of size of input.

\cup

- Removes duplicates:
- Equivalent to $\delta(R \cup_B S)$

The book is wrong. It states we only need to read S in $M-1$ and do one-tuple-at-a-time for R (page 716)

We can do in one pass if

$$\delta(R \cup_B S) \leq M$$

We can approximate to:

$$\delta(B(R)) + \delta(B(S)) \leq M$$

We can remove duplicates as we read tuples:

if tuple already read, ignore
otherwise \hookrightarrow output
add to read tuples. (13)