

We always read smaller table into M


To compute $R - S$, $R \bowtie S$.

Read S
for every tuple t in R
if t not in S
output
(for — also keep track of those output)

To compute $S - R$, $S \bowtie R$

Read S
for — remove all duplicates at the same time.
For every tuple t in R
if t in S
remove from S
for — remove one duplicate only
Output tuples left in S

Query Execution



Query Compilation

- Parsing. A parse tree is constructed.
 - Create an algebraic expression.
- Query Rewrite:
 - Several equivalent query expression
- Physical plan generation
 - Each expression is converted to an evaluation plan by indicating the alg. to use.



b) and c) are the **query optimizer**
⇒ find best query plan.

- 1) Which algebraic expression is the one leading to the most efficient alg.
- 2) For each operation in the expression which alg. will be used to answer it.
- 3) How should each operation pass data to the next operation.
- 4) How are the relations going to be accessed.

Ex: $R(a, b)$ $S(a, c)$

SELECT * from R natural join S
WHERE $b = c$

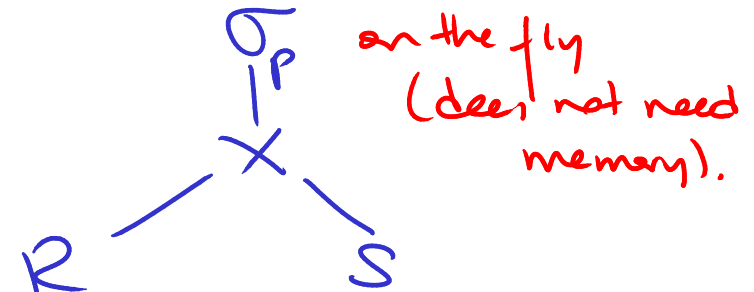
Equivalent Expressions



\bowtie

$$R \bowtie S = \sigma_p (R \times S)$$

Since we can do σ_p on the fly.



But join is common, so DBMS optimize it:

Read S
for every tuple t in R
for every tuple s in S
if t and s satisfy p
output $\text{join}(t, s)$

— B —

Like \cap , \cup , etc. we load smaller table into memory.

But algorithm is different depending on which table is smaller:

$\cap, \cap_B, \times, \bowtie, -, -B$.

- All commutative operations.
- Keep smaller table in memory (plus data structures for fast access).
- Plus at most one block for other table: One pass if, approximately:
 $\min(B(R), B(S)) \leq M$.

Specifically for each of these operations:
 Because they are commutative, assume
 $B(R) \geq B(S)$

\cap, \cap_B

Read S, organize in data structure.

for every tuple t in R

if t in S

if bag op \Rightarrow output t if needed

otherwise output t first time only.

\times

Read S

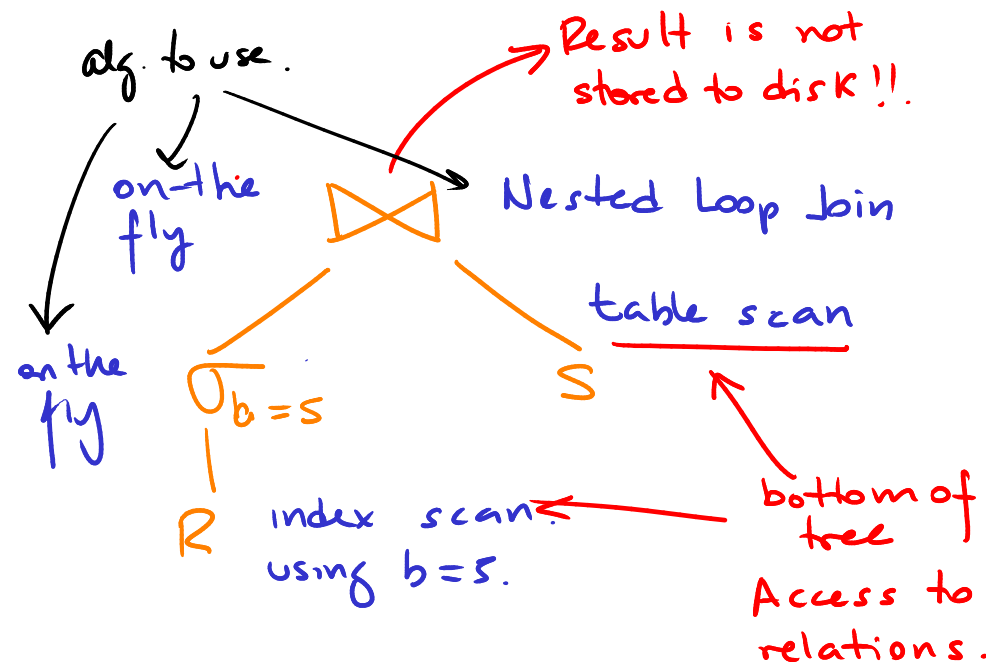
for every tuple t in R

for every tuple s in S

compute cross product, output.

(14)

Annotate tree with algorithms and access methods.



Estimate cost.

\Rightarrow choose fastest!

Access to tuples:

- Sequential scan of heap of Rel.

or

- Using an index to scan a subset of tuples of R (index scan)

Result of query:

- Kept in memory.

(3)

Iterators:

- Many operations access only one tuple at a time.
 - read tuple.
 - inspect
 - dispose
 - read next tuple...

Open() — initiates the process

GetNext() — return next tuple

close() — ends process

Example:

$\pi_a \sigma_{b=3} R$

π_a on the fly

|

$\sigma_{b=3}$ on the fly

|

seq scan of R

R

π and σ can be implemented as iterators

σ inspects one tuple at a time, sends one tuple at a time to π

No need to store any tuple in memory

④

U_B

- Similar to π :
 - We only need to inspect one tuple at a time.
 - $M = 1$. regardless of size of input.

U

- Removes duplicates:
- Equivalent to $\delta(R \cup_B S)$

The book is wrong. It states we only need to read S in $M-1$ and do one-tuple-at-a-time for R (page 716)

We can do in one pass if

$$\delta(R \cup_B S) \leq M$$

We can approximate to:

$$\delta(B(R)) + \delta(B(S)) \leq M$$

We can remove duplicates as we read tuples:

if tuple already read, ignore
otherwise \hookrightarrow output
add to read tuples. ⑬

One Pass alg. for binary operations:

$\cup, \cap, -, \times, \bowtie$

In practice set operations of two types:

- The sets: No duplicates (default).
- Bags: duplicates.

UNION
INTERSECT
EXCEPT

} ALL

\Rightarrow Represented $\cup_B, \cap_B, -_B$

TABLE R UNION ALL TABLE S

Result contains all tuples in R plus all tuples in S.

TABLE R INTERSECT ALL TABLE S

if a tuple in R has m duplicates in R
and n duplicates in S
result contains $\min(m, n)$ duplicates
of tuple.

TABLE R EXCEPT ALL TABLE S

if a tuple in R has m duplicates in R
and n duplicates in S
result contains $\min(m - n, 0)$

(12)

Parameters to measure cost

M. Amount of memory available
in number of blocks

$B(R)$ # of blocks used by heap of R

$|R|$ # of tuples of R (book uses
 $T(R)$)

$V(R, a)$ # of different values of att a
in R

In general:

$V(R, [a_1, a_2 \dots a_n])$

$= |\gamma_{a_1, a_2 \dots a_n} R|$

\Rightarrow # of different values for tuple
 $a_1 \dots a_n$

Cost Model

- We assume that the major component of cost is I/O
- Cost of read equal to cost of write
- Cost of random access of pages equal to cost of seq access.

(5)

Algorithms to answer queries.

2 main classifications.

a) based on type of algorithm:

- 1) Sorting based
- 2) Hash based
- 3) Index based

b) based on difficulty.

- 1) One-pass: Relations are read only once.
- 2) Two passers.
 - Read data (1st pass)
 - Process.
 - Write data.
 - Read data again. (2nd pass).
- 3) Three or more passers.
(needed for very large relations).

2nd pass might read diff number of blocks than 1st pass.

- 3) Three or more passers.
(needed for very large relations).
- Generalization of Two passers.

We can do it in one pass if we have enough memory to

- hold all different groups
- data structures for quick access to groups.
- any data required to compute grouping function.

In general size of tuple of result much smaller than original tuple.

So we simplify

We can do group-by in one pass if

$$B(R) \leq M$$

- $\min(x)$ / $\max(x)$ } Keep current min/max
- $\text{sum}(x)$: Keep current sum
- $\text{count}(x)$ Keep current count
- $\text{avg}(x)$ Keep both current count and sum.

We cannot output tuples until we have read all input tuples.

- We must also create access structures in memory (hash tables, b+ trees) to efficiently find group tuple belongs to.


• In general

- The amount of memory required per group is small.
- Proportional to the number of different groups.

$$|\{ \langle a_1, \dots, a_i \rangle \mid \exists t \in R, t \text{ has attributes } a_1, \dots, a_i \}| \propto V(R, a_1, \dots, a_i)$$

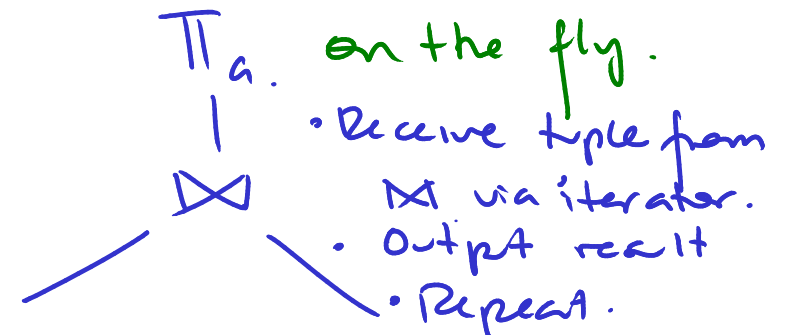
One Pass Alg.

- 1) Tuple-at-a time Π, σ
 - We can read one block at a time.
 \Rightarrow use one memory buffer.



or
if we received tuples from another operation, one tuple at a time with no need for buffering.

(on the fly — no memory needed)



No block in memory needed.

But assume 1 block for simplicity's sake.

Other one pass unary operators.

Duplicate elimination (δ)


• Read each tuple.

• If we have seen it, ignore

• Otherwise output and keep track of it.

We need to keep a copy of each distinct tuple.

input
tuples
(iterator or
from R heap)



We do not need block for output.

⇒ tuples in result output immediately.

We can do δR in one pass
as long as:

$$B(\delta(R)) \leq M-1.$$

Book uses:

$$B(\delta(R)) \leq M$$

because $M \gg 1$

So use latter for consistency.

But, how do we know $B(\delta(R))$
without calculating $\delta(R)$ first?

⇒ Statr.

$R(a_1, a_2 \dots a_n)$

then.

We can use $V(R, a_1 \dots a_n)$ and
the size of the tuple in R to
calculate $\delta(R)$.

Group By:

Generalization of $\delta(R)$

Remember

$$\delta(R) = \gamma^{a_1 \dots a_n} R$$

For $\gamma^{<attlist>} R$.
 $<explist>$

We need to keep track of:

- Each different value of $<attlist>$.
- Info needed to compute $<explist>$.