

SQL and transactions

When connecting to the DBMS operations might be autocommit: each SQL statement is a transaction. (default in psql).

Otherwise transactions start with first SQL statement.

- End with $\begin{cases} \text{COMMIT;} \\ \text{ROLLBACK;} \end{cases}$ — aborts trans.

Read Only Transactions.

Read only transactions can never create conflicts to other transactions.

⇒ easier to interleave.

When possible, indicate transaction is read only:

`SET TRANSACTION READ ONLY;`

Must be first statement of transaction.

By default transactions are read/write

Isolation Levels

Sometimes we are willing to sacrifice serializability for the sake of performance. SQL gives 4 options called **isolation levels**.

Isolation level	Dirty Reads	Non Repeatable Reads	Phantoms
Read Uncommitted	✓	✓	✓
Read Committed	X	✓	✓
Repeatable Read	X	X	✓
Serializable	X	X	X

Type of conflict that the isolation level allows (✓) or not (X)

Use:

SET TRANSACTION ISOLATION LEVEL

<isolation level>;

Must be 1st statement of transaction.

Isolation Levels

We will make the following assumption.

- SQL statements are atomic

READ UNCOMMITTED

- transactions see immediately the effects of other transaction.

Assume R

a	b
1	∅

<p>T₁ BEGIN UPDATE R set b = 5 where a = 1; abort;</p>	<p>T₂ BEGIN READ UNCOMMITTED select * from R ⇒ <table border="1"><tr><td>1</td><td>5</td></tr></table></p>	1	5
1	5		

Changes are immediately visible to T₂ before T₁ ends.

its Iso.
Level
does not
matter


Non Serializable!

READ COMMITTED

- Transactions see effects of other T_i after T_i has committed.
- As if every T has its own local copy of tuple.
- When a T commits, all local copies updated

Example

	a	b
R	1	∅



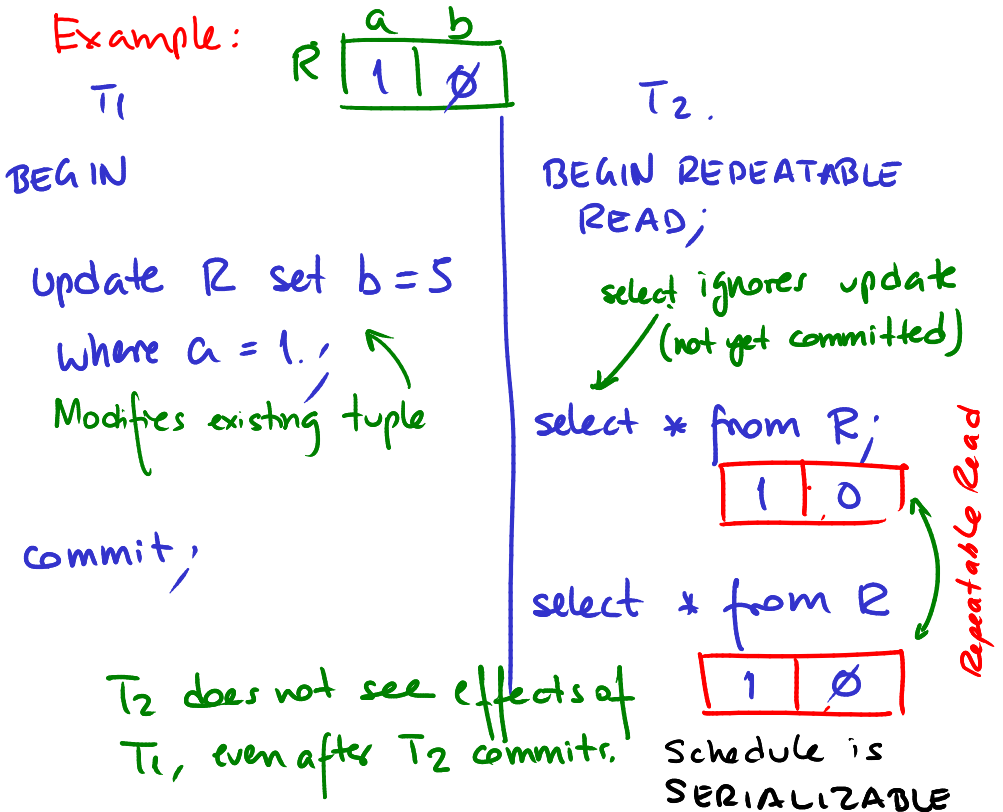
Results in Unrepeatable Read.

Again NON SERIALIZABLE SCHEDULE

REPEATABLE READ

- Look as if the T has a local copy of all tuples it accesses.
- Once it reads a tuple it does not see the effects of other transactions on those tuples. (guarantees repeatable read)
- But it might immediately see inserted tuples by other transactions that have been committed (phantoms).

Example:



Another example

	a	b
R	1	\emptyset

T_1

BEGIN

UPDATE R SET b = 5
where a = 1;

Insert into R values
(0, 0); Phantom.
commit;

T_2 sees the
phantom but
not updated
tuple.

If T_2 is
serializable, T_2 would not see phantom.

T_2

BEGIN REPEATABLE READ

SELECT * from R

\Rightarrow

1	\emptyset
---	-------------

 It does not
see effect of
 T_1
(no dirty read)

SELECT * from R

\Rightarrow

1	\emptyset
\emptyset	\emptyset

\Rightarrow Non Serializable.
schedule.

WHAT IF T_2 IS READ/WRITE?

- All previous examples T_2 is Read Only.
- If two transactions write to the same object one T waits for the other to commit regardless of isolation level!!
- Can result in deadlock even if READ UNCOMMITTED
- If two Tran. write diff. objects, they proceed as previously described.

EXAMPLE


Both T are READ UNCOMMITTED

S_1 = Update R set $b = 5$ where $a = 1$;

S_2 = Update R set $b = \emptyset$ where $a = 2$

$T_1 = S_1 S_2$

$T_2 = S_2 S_1$



DEADLOCK!! one gets killed,
the other continues

- In practice, inserts are allowed to continue
 - No locking of entire table
 - But Transaction might be aborted if T is serializable (and schedule no longer possible to serialize).
- Affects potential interleaving.

EXAMPLE: T_1, T_2 SERIALIZABLE

T_1 :

insert into R(a)
select cant(*)
from R;

T_1 has not committed
so T_2 does not see
tuple.

commit

T_2 :

insert into R(a)
select cant(x)
from R

creates phantom!!

T_2 must be aborted

NO LONGER

serializable.